Workshop Data Platform
Data Import Handbook

Fictional document created for the InterSystems IRIS Vector Search workshop. Version 2026.11

1. Importing data from CSV files

To import a CSV file, open Data, then Imports, and choose "New import". Select the target table, upload the file, and map each CSV column to a table column. The first row must contain column headers.

Files must be UTF-8 encoded and use a comma or semicolon as the separator. Dates must use the ISO format YYYY-MM-DD. The import preview shows the first 20 rows so you can check the column mapping before you start.

When the import finishes, the Imports page shows the number of rows loaded and the number of rows rejected. Rejected rows are listed with the reason, such as a missing required value or a date in the wrong format.

2. Scheduling recurring imports

A recurring import loads files from a watched folder on a schedule. Create one from an existing import by choosing "Repeat on schedule" and selecting hourly, daily or weekly runs.

Each run picks up every new file in the watched folder that matches the file pattern, for example orders_*.csv. Files that were already imported are skipped based on their name and checksum.

Recurring imports run under the account of the user who created them. If that account is disabled, the schedule pauses until another user takes ownership of the import.

3. Import limits and supported formats

A single import file can be up to 2 GB. Supported formats are CSV, JSON Lines and Parquet. Compressed files in gzip format are accepted and are decompressed automatically.

Each workspace can run at most three imports at the same time. Additional imports wait in a queue and start in the order in which they were submitted.

Very wide tables with more than 500 columns cannot be imported through the web interface. Use the bulk loading API for those tables instead.